

Project 3-Intro To xv6 Virtual Memory

Objectives

1. Modify memory layout to move stack to top of address space
2. Implement stack growth

This project can be done with a single partner.

Part 1: Changing memory layout (70%)

In this part, you'll be making changes to the xv6 memory layout. Sound simple? Well, there are a few tricky details.

Details

In xv6, the VM system uses a simple two-level page table. If you do not remember the details, read [OSTEP Section 20.3](#). However, you may find the description in [Chapter 1&2 of the xv6 book](#) sufficient (and more relevant to the assignment).

The xv6 address space is currently set up like this:

```
code
guard page (inaccessible, one page)
stack (fixed-sized, one page)
heap (grows towards the high-end of the address space)
```

In this part of the xv6 project, you'll rearrange the address space to look more like Linux:

```
code
heap (grows towards the high-end of the address space)
... (gap)
stack (at end of address space; grows backwards)
```

This will take a little work on your part. First, you'll have to figure out where xv6 allocates and initializes the user stack; then, you'll have to figure out how to change that to use a page at the high-end of the xv6 user address space, instead of one between the code and heap.

You can see the general map of the kernel memory in `memlayout.h`; the user memory starts at 0 and goes up to `KERNBASE`(exclusive). Thus your stack page

should live at $(\text{KERNBASE} - \text{PGSIZE}) \sim (\text{KERNBASE} - 1)$. **Note that we will not be changing the kernel memory layout at all, only the user memory layout.**

Right now, the program memory map is determined by how we load the program into memory and set up the page table (so that they are pointing to the right physical pages). This is all implemented in `exec.c` as part of the `exec` system call using the underlying support provided to implement virtual memory in `vm.c`. To change the memory layout, you have to change the `exec` code to load the program and allocate the stack in the new way that we want.

Moving the stack up will give us space to allow it to grow, but it complicates a few things. For example, right now xv6 keeps track of the end of the virtual address space using one value (`sz`). Now you have to keep more information potentially e.g., the end of the bottom part of the user memory (i.e., the top of the heap, which is called `brk` in `un*x`), and bottom page of the stack.

Once you figure out in `exec.c` where xv6 allocates and initializes the user stack; then, you'll have to figure out how to change that to use a page at the high-end of the xv6 user address space, instead of one between the code and heap.

Some tricky parts: Let me re-emphasize: one thing you'll have to be very careful with is how xv6 currently tracks the size of a process's address space (currently with the `sz` field in the `proc` struct). There are a number of places in the code where this is used (e.g., to check whether an argument passed into the kernel is valid; to copy the address space). We recommend keeping this field to track the size of the code and heap, but doing some other accounting to track the stack, and changing all relevant code (i.e., that used to deal with `sz`) to now work with your new accounting. Note that this potentially includes the memory growing code that you are writing for part 2.

Part 2: Growing the Stack (30%)

The final item, which is challenging: automatically growing the stack backwards when needed. Getting this to work will make you into a kernel boss. Briefly, here is what you need to do. When the stack grows beyond its allocated page(s) it will cause a page fault because it is accessing an unmapped page. If you look in `traps.h`, this trap is `T_PGFLT` which is currently not handled in our trap handler in `trap.c`. This means that it goes to the default handling of unknown traps, and causes a kernel panic.

So, the first step is to add a case in trap to handle page faults. For now, your trap handler should simply check if the page fault was caused by an access to **the page**

right under the current top of the stack. If this is the case, we allocate and map the page, and we are done. If the page fault is caused by a different address, we can go to the default handler and do a kernel panic like we did before.

Tester:

The testing program `stacktester.c` is shown below and provide in the starter code. Read through the comments to understand the three test cases.

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fs.h"
#include "memlayout.h"
#include "mmu.h"

int magic_num = 123456;

// this address is stored in data/text section
int* stack_ptr = (int*) KERNBASE - 4;

int main(int argc, char *argv[]) {
    // create two integers, one in stack, another in heap
    int stack_val = magic_num;
    int* heap_ptr = malloc(sizeof(int));
    printf(1, "addr of stack_val: %x\n", &stack_val);
    printf(1, "addr of heap_ptr: %x\n", heap_ptr);

    // compare the addresses of the two integers
    // test will pass if stack_val has higher address
    if (heap_ptr < &stack_val) {
        printf(1, "TEST 1 PASSED\n");
    } else {
        exit();
    }

    // initialize somewhere in the end of user memory(i.e., stack)
    *stack_ptr = magic_num;

    // create a child with the same address space
    int rc = fork();
```

```

if (rc == 0) {
    // test would pass only if the stack is properly copied
    // if not, the kernel traps
    if (*stack_ptr == magic_num) {
        printf(1, "TEST 2 PASSED\n");
    }
} else if (rc > 0) {
    (void) wait();
    // grow the stack by accessing a address
    // in the page right below the initial stack page
    char *p = (char*) (KERNBASE - PGSIZE - 4);
    *p = 'z';

    // test would pass if trap is handled properly
    printf(1, "TEST 3 PASSED\n");
} else {
    printf(1, "fork failed!\n");
}
exit();
}

```

The testing program should produce something similar to the following output. The **stacktester** executable (when booted into xv6) is based on the program above. First you should copy **stacktester.c** in xv6 folder, and modify the **Makefile** to add **stacktester** as a new user program.

```

$ stacktester
addr of stack_val: 7FFFFFFAC
addr of heap_ptr: 8FF8
TEST 1 PASSED
TEST 2 PASSED
TEST 3 PASSED

```

Bonus question (10%)

Write code to try and get the stack to grow into the heap. Were you able to? If not explain why in detail showing the relevant code and output.

Hints

Rearrange the stack

Specifically, let's start by opening up `exec.c` check the `exec` function which implements the system call. Exec does the following:

Part 1. Opens the executable file and parses it. Typically executable files are made up of a header including information that allows us to index the rest of the file. The file after that consists of sections including the code, the global data, and sometimes other sections like uninitialized data. These are the parts of the memory that we need to initialize from the executable. The header information includes the number of sections, the start of each section in the file, and where it maps to in virtual memory, and the length of each section.

Part 2. Initializes the kernel memory using `setupkvm()` which maps the pages of the kernel to the process address space. We don't really need to know what happens in here.

Part 3. It then moves on to load the sections of the executable file into memory using `loadvm()` which creates the memory pages for each section and maps them to the address space (by initializing the page table pointers -- more details in a bit). These sections in xv6 are loaded starting at VA 0, and going up. Each new section starts at the beginning of a new page. Recall that sections include code, global/static data, etc. Conveniently, we can then keep track of where the user address space ends, which also defines the size of the process using one value (`proc->sz`). So, as we map new pages, `sz` (rounded up to the next page) can serve as their virtual address since we are simply filling in the address space sequentially.

In VM, we map virtual pages to physical frames. XV6 has no swap so all memory pages have to be in physical memory. Physical memory is allocated by the kernel allocator `kalloc`. If you want to dig deeper you will see how these pages are initialized using `kinit()` and `kinit2()` which are called during the boot process in `main.c`. As a result, we use `kalloc()` as we request each new page inside `vm.c` to allocate a new physical page. The `vm.c` functions such as `allocvm()` typically follow this up with a call to `mappages` which is used to initialize the page table entries mapping the virtual address to the physical page that it just allocated. Otherwise these physical frames that we allocate cannot be used by our process.

Part 4. At this point, we loaded code and data sections, and it's time to create the stack. xv6 does not support a heap at the moment (there is no `malloc()/free()`)

available to user programs if you noticed). It currently maps the stack in its virtual address space at a page right after the last page we loaded from the executable (i.e., at `sz` rounded up to the next page boundary). Since the stack grows down, allocating a page here means there is no room to grow the stack -- as it grows down, it will run into the code/data. To protect against that, xv6 adds one page buffer and marks it as inaccessible so that in case the stack grows, we get a memory error and can stop the program. The code to create the stack is:

```
sz = PGROUNDUP(sz); //round sz up to the next page boundary since stack  
must start in a new page  
if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0) // our first  
introduction to allocuvm; it allocates and maps two pages goto bad;  
clearpteu(pgdir, (char*)(sz - 2*PGSIZE)); //we clear the PTE for the first  
page to create a buffer page between stack and code/data
```

TODO 1: This is the part of the code that we need to change to move the stack. The current code calls `allocuvm` to create two pages, one for the stack and one as a guard page starting at VA `sz` which is right after the code and data. It then clears the page table entry for the guard page. We want to locate the stack starting at the top of the user address space to give it room to grow. To understand what we need to do, let's look at `allocuvm`. It takes 3 parameters: a. The page table (`pgdir`). This will not change b. The virtual address of the first page we are mapping -- this needs to change to point to the top page of the user part of memory (right under `KERNBASE`). If you use `KERNBASE`, you will try to map the page to the kernel address space. c. The virtual address of the last page we are mapping. For us, we are creating a stack with only a single page, so this can another address in the same page, slightly bigger than the first address. `Allocuvm` allocates the page, and maps it to the page table. So, basically we are done with moving the stack by just changing these parameters to the right value. However, there are a few loose ends to tie up.

Part 5: Finally, we initialize the stack pointer, currently to `sz`.

TODO 2: you will have to change this to the address of the top word in the stack page. Note that `KERNBASE` is the first word in the kernel address space, so this is the word right under that. We proceed to initialize the stack with the parameters for `main` as per the linux/x86 stack convention. The details are not important to us for now.

Loose Ends/Other changes. Now that we moved the stack, a few places in the Kernel that hard coded the previous location of the stack have to be changed. These include:

TODO 3: All of the functions that are defined in `syscall.c` (and `sysfile.c`) for accessing the user stack have some checks to see if the addresses are indeed on the stack. These checks compare the address against `sz` since that was the top of the stack in the old implementation. You have to change those checks (or remove them if it is easier). Check all the accessor functions such as `argint`, `argstr`, `argptr`, `argfd`, etc.

TODO 4: `copyvm()`. This function is used as part of `fork()` to create a copy of the address space of the parent process calling fork to the child process. It is implemented in `vm.c`. If you look at this function, it is one big for loop that iterates over the virtual address space and copies the pages one by one. The loop starts with:

```
for(i = 0; i < sz; i += PGSIZE)
```

since it assumes the virtual address space starts at 0 and goes to `sz`. Now this has to be changed to take into account the new stack. If we look deeper, it reads the page table to get the PTE for the page, allocates a new physical frames, and copies the page from the parent memory to the new page. Finally it uses `mappages` to map this new copy to the child address space by adding a PTE to its page table. How do we change it? Now `sz` tracks the bottom part of the address space, so its ok to leave that loop alone. We have to keep track of the size of the stack, and added another loop that iterates over the stack page(s) and does the same thing (`kalloc` a page for each one, `memmove` to create a copy from the parent, and then `mappages()` to add it to page table). The loop will be very similar, with the exception of the virtual address ranges that iterates over. Before we add stack growth, the stack is only one page, but as the stack grows we need to keep track of the number of stack pages. To prepare for this, we need to add a variable in struct `proc` to keep track of the size of the stack (in pages or bytes--either is fine, but I recommend pages). This counter starts out with a stack of one page; set it in `exec()`.

Debugging: If your stack moved correctly, xv6 will be able to boot into shell successfully. If you don't allocate/map the stack correctly, you will get errors either in the allocation functions (e.g., remapping errors) or as your program runs (page faults). If you don't take care of all the `argint()` etc. functions some of your system calls will not be able to pass parameters correctly. The results could be weird. For example, `printf` won't print, and `wait` won't wait (leading to `init` continuing to fork processes, etc...)

Growing the stack

Now that our stack has been moved, we have room to grow it. When the a program causes the stack to grow to an offset bigger than one page, at this point, we will be accessing a page that is not allocated/mapped. This will cause a page fault. Basically, we will trap to the trap handler in `trap.c`. In there there is a switch statement with a case for every supported trap. We need to add a case for page faults. This page fault has trap number 14 (or `T_PGFLT`) as defined in `traps.h`.

TODO 5: Add a case for the page fault. When a page fault occurs, you can check the address that caused the page fault in a hardware register called CR2. The CR register (standing for Control Register) keep track of important hardware state information. You can read the CR2 register using the function `rcr2()`. Once you have the offending address, next we need to check if it is from the page right under the current bottom of the stack. If it is, we need to grow the stack. You can use `allocvm` again, but you have to initialize it with the right parameters to allocate one page at the right place. After that, you can increment your stack size counter, which finishes your trap handler.

Voila! you should be good to go.

Submission

Upload all of your source files (but not `.o` files, please, or binaries!) as a zip file `xv6-public.zip` to Blackboard. A simple way to do this is type `make` in your `xv6-public` folder to make sure it builds, and then type `make clean` to remove unneeded files. Generate a `patch` for you kernel and upload it. Also, please make a **SCREENSHOT** file, which shows your output of `stacktester`. Finally, please indicate who is your partner on Blackboard.